

Image Classification with CNN: Cats vs Dogs

...

Jack D

Disclaimer

This code was developed via a guide published by Jay 'Coding Lane' Patel on YouTube.

The following slides regard the explanation of how CNN's work and the python implementation of the neural network.

Intro to Data

- Kaggle's Cats and Dogs dataset
 - Not all is used (a lot of data)
 - Original Dataset has approx 12,000 pictures
 - We will only use 2,400
 - <https://www.microsoft.com/en-us/download/details.aspx?id=54765>
- Data transformed externally from JPG to CSV (containing RGB data)
- Each image is 100*100 pixels (each containing 3 rgb values)

ex:



ex:



ex:



ex:



ex:



Preprocessing

- Original dataset shape is (2,400 by 30,000)
- Needed to reshape to (2,400 tuples of 100*100*3(rgb))

```
xtrain.shape,ytrain.shape,xtest.shape,ytest.shape
```

```
((2000, 30000), (2000,), (400, 30000), (400,))
```

```
xtrain.shape,ytrain.shape,xtest.shape,ytest.shape
```

```
((2000, 100, 100, 3), (2000, 1), (400, 100, 100, 3), (400, 1))
```

```
xtrain.reshape(len(xtrain), 100, 100, 3)
```

```
ytrain.reshape(len(ytrain), 1)
```

```
xtest.reshape(len(xtest), 100, 100, 3)
```


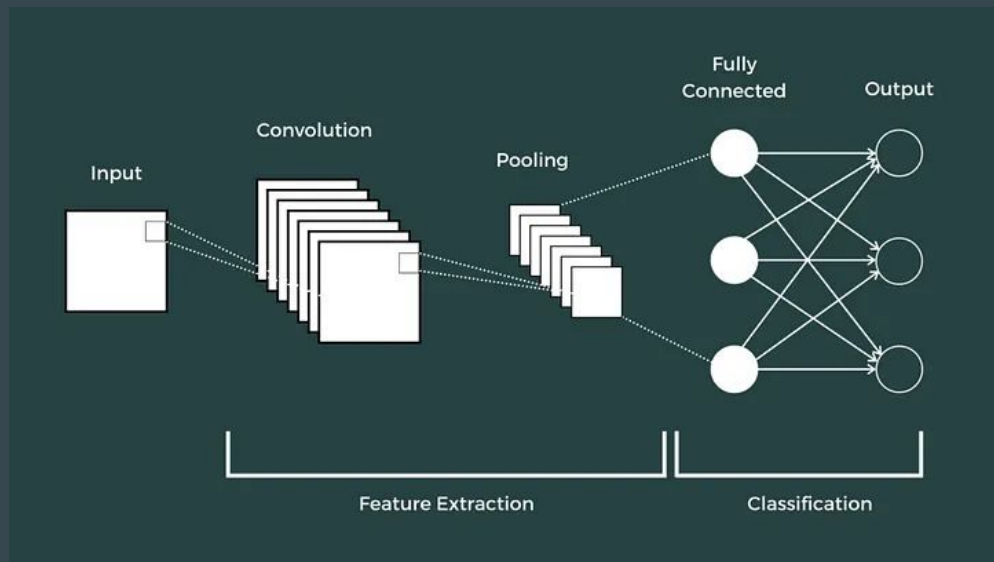
```
ytest.reshape(len(ytest), 1)
```

- Need to rescale RGB values from 0-255 to 0-1 for input

```
xtrain = xtrain/255.0
```

```
xtest = xtest/255.0 #normalize pixels w/ numpy
```

Intro: Convolutional Neural Networks (CNN)



1	0	2	3
4	6	6	8
3	1	1	0
1	2	2	4

→

6	8
3	4

- Convolution - aggregates partitions of original input
- Pooling - amplifies features of image to increase distinctiveness
 - These, and some ReLU nodes, make up our 'hidden layer' in the network

Python Implementation of CNN

```
model = Sequential([  
    #convolution  
    Conv2D(32, (3,3), activation = 'relu', input_shape = (100,100,3)),  
    MaxPooling2D((2,2)), #default value 2  
    Conv2D(32, (3,3), activation = 'relu'),  
    MaxPooling2D((2,2)),  
  
    #neural net  
    Flatten(), #turn into inputs for NN  
    Dense(64, activation = 'relu'), #fully connected layer of nodes  
    Dense(1, activation = 'sigmoid') #binary classification  
)
```

Compiler:

- Minimizes binary entropy.
- Uses adam optimizer (faster training than SGD, developed by OpenAI).
- Assesses Accuracy

Model:

- Convolution of Image = Conv2D (32 filters, 3x3), uses ReLU (ie 0 for activation)
- Pooling = MaxPooling (2x2, steps 2)
- Flatten = convert image to interpretable array
- Dense = make layer of nodes in NN which will connect to all previous and consecutive nodes

```
#opt = keras.optimizers.SGD(learning_rate=0.01) #specify learning rate, hyperparam  
model.compile(loss = 'binary_crossentropy', optimizer = 'adam', metrics = ['accuracy'])
```

```
model.fit(xtrain, ytrain, epochs = 5, batch_size = 64)
```

```
Epoch 1/5  
32/32 ————— 2s 68ms/step - accuracy: 0.7053 - loss: 0.5309  
Epoch 2/5  
32/32 ————— 2s 68ms/step - accuracy: 0.7404 - loss: 0.4773  
Epoch 3/5  
32/32 ————— 2s 67ms/step - accuracy: 0.8023 - loss: 0.4202  
Epoch 4/5  
32/32 ————— 2s 66ms/step - accuracy: 0.8240 - loss: 0.3699  
Epoch 5/5  
32/32 ————— 2s 67ms/step - accuracy: 0.8536 - loss: 0.3015
```

```
<keras.src.callbacks.history.History at 0x1df19af2760>
```

Python Implementation of CNN

Training of Model and Evaluation

Hyperparameters Examined

- Different numbers of epochs (training iterations)
- Stochastic Gradient Descent vs. Adam optimization

```
model.fit(xtrain, ytrain, epochs = 5, batch_size = 64)
```

Epoch 1/5

32/32 ————— 2s 68ms/step - accuracy: 0.7053 - loss: 0.5309

Epoch 2/5

32/32 ————— 2s 68ms/step - accuracy: 0.7404 - loss: 0.4773

Epoch 3/5

32/32 ————— 2s 67ms/step - accuracy: 0.8023 - loss: 0.4202

Epoch 4/5

32/32 ————— 2s 66ms/step - accuracy: 0.8240 - loss: 0.3699

Epoch 5/5

32/32 ————— 2s 67ms/step - accuracy: 0.8536 - loss: 0.3015

<keras.src.callbacks.history.History at 0x1df19af2760>

```
model.evaluate(xtest, ytest)
```

13/13 ————— 0s 21ms/step - accuracy: 0.2058 - loss: 0.6967

[0.6933041214942932, 0.4950000047683716]

Model Evaluation - SGD($\alpha=0.001$)

- Epoch = 1
 - [loss, acc]

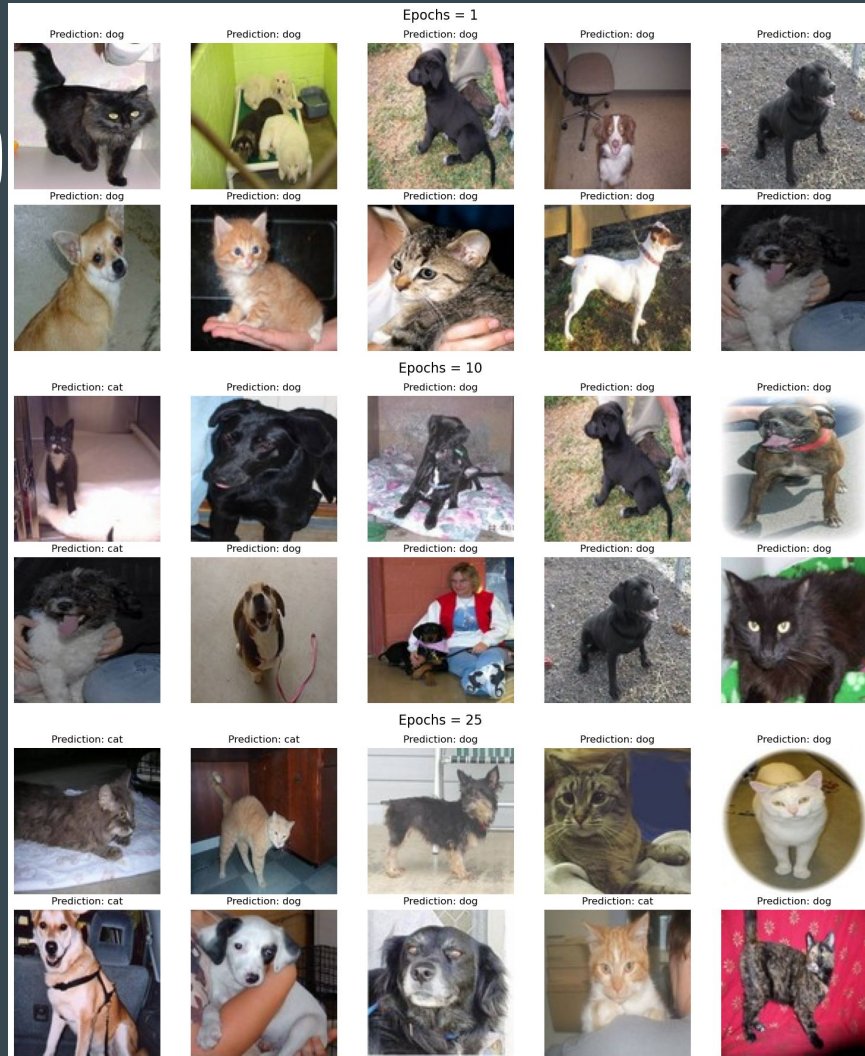
```
[0.7023858428001404, 0.5]
```

- Epoch = 10

```
[0.6880124807357788, 0.5350000262260437]
```

- Epoch = 25

```
[0.6830646395683289, 0.5849999785423279]
```



Model Evaluation - Adam

- Epoch = 1
 - [loss, acc]

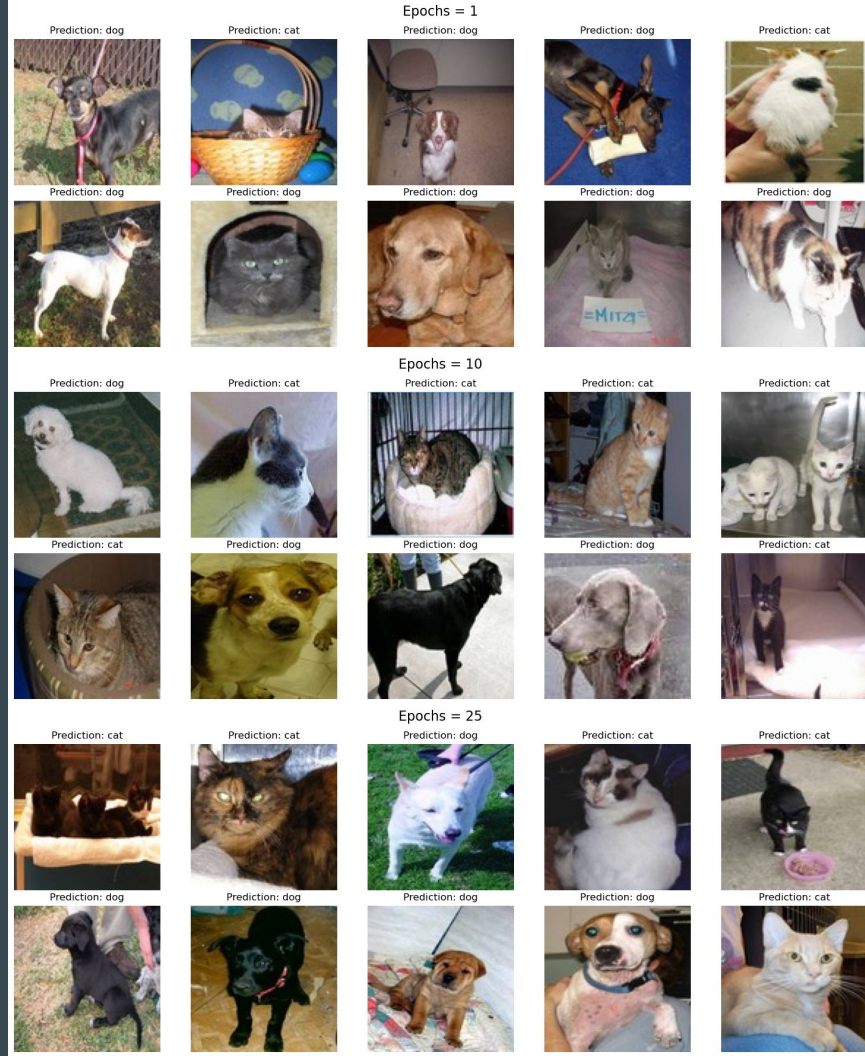
[0.6836750507354736, 0.5425000190734863]

- Epoch = 10

[0.7136306166648865, 0.6650000214576721]

- Epoch = 25

[1.615975022315979, 0.7024999856948853]



Conclusion

- The more training iterations (epochs) for a neural network, the higher the accuracy
- In the context of CNN, Adam optimizers were able to utilize epochs more than SGD
- CNN plays a fundamental role in computer vision